


Trust & Safety Operations: SLA Recovery Briefing



Timing: 0:00 – 1:30 (90s) • Pace: Confident, direct

Data Source: dashboard-ready.csv (2,800 reviews, Sep 6 – Oct 5, 2026)

Core Talking Points

1. The headline: decision quality is strong (98.6% average audited QA score; 333 of 340 audits error-free), but SLA attainment is 78.2% against a 95.0% contractual target.
2. The discovery: the deficit is not systemic. Commercial Spam and Fraud & Scams together run at 96.7% SLA.
3. The concentration: 565 of 610 SLA breaches (93%) sit in Violence_HighRisk, Hate_Speech and Emergency_TCO. A further 209 breaches (34% of all breaches) landed in just 4 days, Sep 24–27, when the intake mix shifted toward high-risk content and SLA fell to 41.1%.
4. The takeaway: volume did not grow, so this is a routing and capacity-allocation problem, not a headcount problem. The target is ≥95% SLA within 14 operational days with no new hires.

Spoken Script

Good morning, leadership team. My name is Wesley Taguinod, and today I am presenting the operational performance briefing and SLA recovery strategy for our Content Moderation and Trust & Safety operations.

[Pause 1s — glance at audience]

Across 2,800 reviews over the past 30 days, our account operates at:

- A 98.6% average audited QA score,
- A weighted average handling time of 52.5 seconds, and
- An overall SLA attainment of 78.2% — trailing our contractual 95.0% commitment.

[Gesture toward right KPI cards]

Here is the bottom line up front: our challenge is not a quality failure, and it is not widespread. Of 340 audited decisions, 333 were error-free. Our false negative rate — harmful content mistakenly left live — is just 0.59%.

Our high-volume deterministic queues, Commercial Spam and Fraud, are running above target at 96.7% SLA.

The breach is concentrated in two places:

1. Our contextual and high-risk queues — Violence, Hate Speech and Emergency TCO — hold 93% of all SLA breaches.
2. A 4-day window, September 24 to 27, when the intake mix shifted sharply toward high-risk video and SLA fell to 41%.

Daily volume did not rise during that window. That tells us this is an allocation problem, not a headcount problem. By shifting existing capacity toward the high-risk queues when the mix changes, tightening how we route video, and calibrating our newest agents, we are targeting 95% SLA or better within 14 operational days. Let me show you what the data reveals.

[Transition: advance to Slide 2]

Baseline Operations: Queue & Modality Breakdown

WHAT? • BASELINE OPERATIONS

Baseline Operations: Queue & Modality Breakdown


2,800 reviews

15 active agents

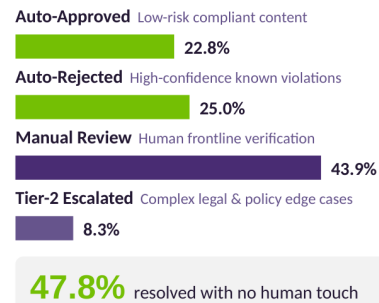
5 queues

4 media modalities

Queue SLA attainment vs. 95.0% contractual target



Human-in-the-Loop triage funnel



ApexTrust Operations • Content Moderation Executive Briefing • Sep 6 – Oct 5, 2026

2

Timing: 1:30 – 3:30 (120s) • Pace: Methodical, grounded

Data Source: dashboard-ready.csv (2,800 reviews, Sep 6 – Oct 5, 2026)

Core Talking Points

1. Mathematical rigor: AHT is duration-weighted across total handling seconds, and SLA is calculated at the transaction level (610 breached of 2,800), not by averaging queue percentages.
2. The queue divide: Spam & Fraud = 1,346 reviews (48.1% of volume) at 12–36s AHT and 96.7% SLA. Violence, Hate Speech & TCO = 1,454 reviews at 44–110s AHT and 61.1% SLA, holding 565 of 610 breaches (93%).
3. HITL efficiency: automation resolves 47.8% of volume (22.8% auto-approved, 25.0% auto-rejected). But it is uneven: Spam is 75% automated, Violence 33% and Emergency TCO only 24%.

Spoken Script

To understand where the variance comes from, we broke our 2,800 reviews down across 5 queues and 4 content modalities.

[Point to left chart: Queue SLA Attainment]

The queue-level data shows a sharp split:

- Commercial Spam and Fraud & Scams are model operations. They handle nearly half our volume — 1,346 reviews — at 96.4% and 96.9% SLA, with handling times under 37 seconds.
- Our contextual and high-risk queues pull the account down: Hate Speech at 65.4%, Violence at 57.6% and Emergency TCO at 56.7%. Together these three queues hold 93% of every SLA breach we recorded.

[Point to right visual: HITL Funnel]

Now look at how work reaches agents. Our Human-in-the-Loop automation resolves 47.8% of volume without a human touch — about 23% auto-approved and 25% auto-rejected by hash matching and high-confidence classifiers.

That leaves 44% for frontline manual review and 8% for Tier-2 escalation.

The catch is that automation is uneven: three-quarters of Spam is automated, but only a third of Violence and a quarter of Emergency TCO. The human queues carry the hardest content. To see why they breach, we need to look at media type and what happened over time.

[Transition: advance to Slide 3]

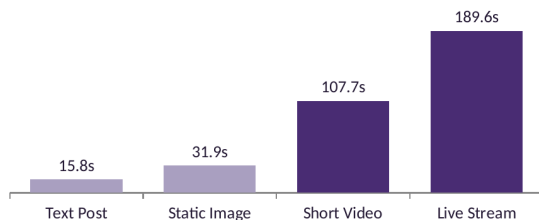
Root Cause: Where the SLA Drag & Backlog Originate

SO WHAT? · ROOT CAUSE

Root Cause: Where the SLA Drag & Backlog Originate



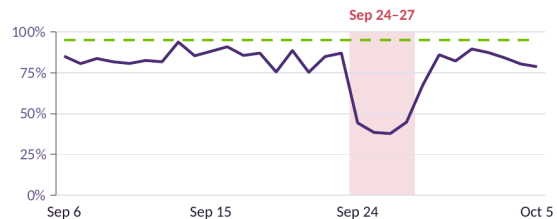
Media multiplier: AHT by content type



Video = 28% of volume, 69% of handling hours · short video 6.8×, live 12× text AHT



Daily SLA attainment vs. 95% target



Volume
95 → 89/day

High-risk mix
24% → 43%

Avg. TAT
134 → 241 min



719 egregious items (26% of volume) sit entirely in Violence, Hate Speech & TCO and trigger 5.5-hour shift caps with 30-minute wellness rotations. Capacity cannot be recovered by extending shifts; it must be recovered by routing.

ApexTrust Operations · Content Moderation Executive Briefing · Sep 6 – Oct 5, 2026

3

Timing: 3:30 – 6:00 (150s) · **Pace:** Analytical core

Data Source: dashboard-ready.csv (2,800 reviews, Sep 6 – Oct 5, 2026)

Core Talking Points

1. The media multiplier: AHT is driven by media type, not agent effort. Text averages 15.8s, static images 31.9s, short video 107.7s (6.8× text) and live streams 189.6s (12× text). Video and live are 28% of volume but 69% of handling hours.
2. The surge window (Sep 24–27): daily volume did not rise (95/day before vs. 89/day during). The mix shifted: Violence + Emergency TCO rose from 24% to 43% of intake, weighted AHT rose 31% (49.7s → 65.3s), average turnaround nearly doubled (134 → 241 min) and SLA fell to 41.1%. Those 4 days produced 209 of 610 breaches.
3. Workforce protection constraint: all 719 egregious items sit in Violence (418), TCO (164) and Hate Speech (137). Exposure caps (5.5-hour shifts) mean the gap cannot be closed by lengthening shifts.

Spoken Script

This brings us to the core diagnostic question: why did SLA fall to 78%, and where did the backlog come from?

The data points to two root causes.

[Focus on left card: Media Modality]

First, media type. Reviewing a text post takes an agent about 16 seconds. A static image takes 32 seconds. Short-form video jumps to 108 seconds, and live streams average over three minutes.

Here is the critical insight: video and live streams are only 28% of our volume, but they consume 69% of all handling hours. Any day the intake tilts toward video, our capacity is stretched.

[Focus on right card: Daily SLA]

Second, look at the daily trend on the right. Most days ran between 75 and 94%. Then, from September 24 to 27, SLA collapsed to roughly 40% a day.

[Point to the shaded window and the three facts below]

Importantly, volume did not spike — we averaged 89 items a day during the window versus 95 before it. What changed was the mix: Violence and Emergency TCO rose from 24% to 43% of our intake. Handling time per item rose 31%, and average turnaround nearly doubled, from 134 to 241 minutes. Emergency TCO, which carries a strict 1-hour legal removal SLA, averaged 78 minutes to decision and met its SLA only 29% of the time in those 4 days.

Furthermore, every one of our 719 egregious items sits in these high-risk queues, and our duty-of-care rules properly cap that exposure at 5.5-hour shifts with wellness rotations.

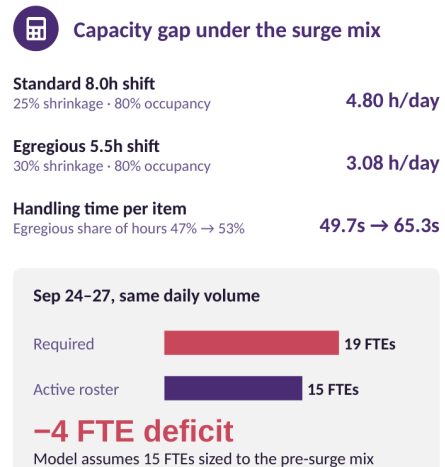
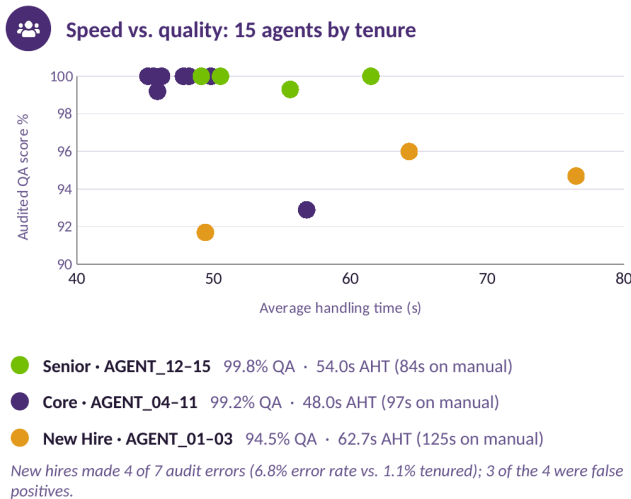
The conclusion: our agents did not slow down. We had a structural mismatch between a video-heavy, high-risk intake mix and a fixed allocation of capacity.

[Transition: advance to Slide 4]

Workforce Dynamics: Speed vs. Quality Matrix

SO WHAT? · WORKFORCE DYNAMICS

Workforce Dynamics: Speed vs. Quality Matrix



ApexTrust Operations · Content Moderation Executive Briefing · Sep 6 – Oct 5, 2026

4

Timing: 6:00 – 8:00 (120s) · Pace: Objective, diagnostic

Data Source: dashboard-ready.csv (2,800 reviews, Sep 6 – Oct 5, 2026)

Core Talking Points

1. Tenure and speed: on manual reviews, where speed matters, seniors average 84.2s vs. 96.8s for core and 125.2s for new hires (33% faster than new hires). Seniors' overall AHT (54.0s) is higher than core (48.0s) only because they carry more egregious work (30% of their items vs. 23%). On short video, core agents are fastest (102s) and new hires slowest (123s, 20% slower); on live streams new hires are 41% slower than core.
2. Error taxonomy: 7 errors in 340 audits (4 false positives, 2 false negatives, 1 wrong policy). New hires made 4 of the 7 (6.8% error rate vs. 1.1% for tenured agents), and 3 of those 4 were false positives on static images — over-enforcement, not missed harm.
3. Capacity math (BA guide Section 5): productive hours = shift × (1 – shrinkage) × occupancy = 4.80h on a standard shift, 3.08h on an egregious 5.5h shift. Applying this to each phase's handling hours, split by egregious vs. standard work: if 15 FTEs covered the pre-surge mix, the Sep 24-27 mix needed 18.9 → 19 FTEs at the same daily volume, a 4-FTE gap. Workload above available servers means Erlang-C queue delay is effectively certain, which matches turnaround nearly doubling.

Spoken Script

To turn this diagnosis into a staffing model, we looked at agent performance and applied the capacity formulas from our operations guide.

[Point to left scatter: Speed vs Quality]

On the left, each dot is one of our 15 agents, plotted by handling time and audited QA score.

The tenure pattern is clear:

- Our senior agents in green — Agents 12 through 15 — average 99.8% QA. Their overall handle time looks higher than core because they take on more of the egregious work, but on manual reviews they are our fastest, at 84 seconds.
- Our core agents in purple deliver 99.2% QA at 48 seconds, and they are our fastest on short video.
- Our three new hires in amber — Agents 01, 02 and 03 — average 94.5% QA and take 125 seconds on manual reviews.

When we look at the errors themselves, new hires account for 4 of the 7 audit errors we found, and 3 of those were false positives — over-enforcement on images, not missed harm.

[Focus on right card: Capacity Gap]

Now the workforce math on the right.

On a standard 8-hour shift, with 25% shrinkage and 80% occupancy, each agent delivers 4.8 productive hours. On egregious queues, the 5.5-hour cap and 30% shrinkage leave just 3.08 productive hours.

During the September 24 to 27 window, the mix shift raised handling time per item from 50 to 65 seconds, and pushed the egregious share of our workload from 47% to 53%. If our 15 agents were sized for the normal mix, that window needed 19 — a 4-agent shortfall, at the same daily volume.

That is the mechanical reason our SLA dipped. The good news: the gap is 4 FTEs for a few days, which is a routing problem we can solve without hiring.

[Transition: advance to Slide 5]

Action Plan: Rebalance Capacity, Not Headcount

NOW WHAT? · ACTION PLAN

Action Plan: Rebalance Capacity, Not Headcount



ACTION 1

Mix-triggered cross-queue flex staffing

- When Violence + TCO exceed 35% of daily intake, flex 4 FTEs from Spam & Fraud
- Spam & Fraud use 22% of handling hours and Spam is already 75% automated



ACTION 2

Smarter video routing & triage pilot

- Route video & live to Core and Senior agents; new hires are 20–41% slower on them
- Pilot wider auto-triage in Violence (33% automated) and TCO (24%) with QA sign-off



ACTION 3

New-hire calibration & wellness scheduling

- Daily 20-minute calibration with senior mentors on over-enforcement (false positives)
- Automate rotation: 5.5h caps and ≤2 consecutive hours in egregious queues

0 new hires

and zero external software spend

78.2% → ≥95%

account SLA target within 14 days

56.7% → ≥95%

TCO 1-hour emergency SLA target

ApexTrust Operations · Content Moderation Executive Briefing · Sep 6 – Oct 5, 2026

Timing: 8:00 – 10:00 (120s) · **Pace:** Decisive, ROI-focused

Data Source: dashboard-ready.csv (2,800 reviews, Sep 6 – Oct 5, 2026)

Core Talking Points

1. Zero budget required: every action reallocates existing capacity, changes routing or adds coaching time; no hires and no new software.
2. Evidence for each action: (1) the surge mix needed 4 more FTEs while Spam & Fraud use only 22% of handling hours for 48% of volume; (2) new hires are 20% slower than core on short video and 41% slower on live, and Violence/TCO automation (33%/24%) trails Spam (75%); (3) new hires made 4 of 7 audit errors, 3 of them false positives.
3. Targets, not forecasts: ≥95% account SLA and ≥95% TCO 1-hour SLA are the contractual targets to hit within 14 days. Track them weekly in the dashboard. Excluding the Sep 24–27 window, SLA was already 83.6%, so most of the gap is the mix-shift response.
4. Sustainable workforce health: exposure caps and rotations stay in place, so speed gains do not come at the cost of burnout or attrition.

Spoken Script

That brings us to the plan: how do we fix this permanently without adding payroll?

We propose three targeted, data-backed interventions.

[Point to Action 1]

Action 1: mix-triggered cross-queue flex staffing.

- Our model shows the high-risk mix needed 4 more FTEs. So when Violence and Emergency TCO exceed 35% of daily intake, we flex 4 agents' worth of capacity out of Spam and Fraud into the high-risk queues.
- Spam and Fraud can absorb this: they use only 22% of our handling hours for nearly half our volume, and three-quarters of Spam is already automated.

[Point to Action 2]

Action 2: smarter video routing and a triage pilot.

- Video and live streams drive 69% of our handling hours. Our new hires are 20% slower than core agents on short video and 41% slower on live streams, so we route that content to core and senior agents.
- We will also pilot wider automated triage in Violence and Emergency TCO, which are only 33% and 24% automated today versus 75% in Spam — with QA sign-off before any threshold change.

[Point to Action 3]

Action 3: new-hire calibration and automated wellness rotation.

- Daily 20-minute calibration huddles pair our 3 new hires with senior mentors, who run at 99.8% QA. The focus is over-enforcement: 3 of the new hires' 4 audit errors were false positives.
- We automate scheduling so no agent spends more than 2 consecutive hours in egregious queues, within the 5.5-hour cap.

[Highlight bottom banner: Recovery Targets]

The executive takeaway:

- This plan needs zero new hires and zero external software spend.
- Our target is to bring account SLA from 78.2% to 95% or better within 14 operational days. Outside the 4-day surge window we were already at 83.6%, so most of the gap is in how we respond to mix shifts.
- Most importantly, we target 95% or better on the Emergency TCO 1-hour legal SLA — up from 56.7% today — protecting platform integrity and our regulatory position.

Thank you for your time. I will now open the floor to questions.

Anticipated Questions & Defensible Responses

APPENDIX · LEADERSHIP Q&A

Anticipated Questions & Defensible Responses



Won't Spam SLA collapse if we flex capacity out of Spam & Fraud?

Unlikely. Spam runs at 12.3s AHT and 96.4% SLA with 75% of items automated, and Spam & Fraud together use only 22% of handling hours. Flexing is triggered only on high-risk mix days.



Why not mandate 8-hour shifts for everyone to clear the backlog?

All 719 egregious items sit in the high-risk queues, and duty-of-care caps exposure at 5.5 hours. Violence QA already dipped to 88.9% during the surge (9 audits). Rotation protects accuracy.



Does widening automated triage risk more false positives?

We pilot with QA sign-off. Today auto-approved items audit at 100% (75 audits) and auto-rejected at 97.3% (81 audits), and appeals overturn only 15.8% of disputed decisions.

BACKUP SLIDE — use during the 2:00 synthesis & leadership Q&A.

Q1: "If we flex capacity from Spam and Fraud into the high-risk queues, won't the Spam SLA collapse?"

"Unlikely. Commercial Spam runs at 12.3 seconds AHT and 96.4% SLA, and 75% of Spam items are resolved automatically. Spam and Fraud together account for 48% of our volume but only 22% of our handling hours. And the flex only triggers on days when Violence and TCO exceed 35% of intake. During the Sep 24–27 window Spam and Fraud also dipped (61% and 68%) because the whole floor was saturated, so relieving the high-risk queues helps every queue."

Q2: "Why not simply mandate 8-hour shifts for all agents to clear the backlog?"

"Because the backlog sits exactly where the exposure is. All 719 egregious items are in Violence, Hate Speech and Emergency TCO, and our duty-of-care policy caps that exposure at 5.5-hour shifts. We already saw a warning sign: audited QA in the Violence queue dipped to 88.9% during the surge window, versus 100% before it (a small sample of 9 audits, but the direction matters). Rotation into non-egregious work protects accuracy without driving attrition."

Q3: "Does widening automated triage increase false positives?"

"That is why it is a pilot with QA sign-off, not a blanket change. Today, automated decisions audit well: auto-approved items scored 100% across 75 audits, and auto-rejected items 97.3% across 81 audits. Disputed decisions also have a safety net: of 190 appeals this period, only 30 (15.8%) were overturned."